

CIFAR-10: VGG-A


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


196/196 [=====] - 8s 39ms/step - loss:
1.4905 - accuracy: 0.9704 - val_loss: 1.7007 - val_accuracy: 0.7588



Run2

Epoch 200/200


196/196 [=====] - 8s 39ms/step - loss:
1.4931 - accuracy: 0.9677 - val_loss: 1.6946 - val_accuracy: 0.7663



Run3

Epoch 200/200

196/196 [=====] - 8s 39ms/step - loss:
1.4912 - accuracy: 0.9697 - val_loss: 1.7026 - val_accuracy: 0.7552




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


196/196 [=====] - 8s 41ms/step - loss:
1.5070 - accuracy: 0.9545 - val_loss: 1.7310 - val_accuracy: 0.7264



Run2

Epoch 200/200


196/196 [=====] - 8s 39ms/step - loss:
1.5091 - accuracy: 0.9527 - val_loss: 1.7407 - val_accuracy: 0.7182



Run3

Epoch 200/200

196/196 [=====] - 8s 40ms/step - loss:
1.5067 - accuracy: 0.9548 - val_loss: 1.7280 - val_accuracy: 0.7328



CIFAR-10: VGG-B


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


196/196 [=====] - 10s 52ms/step - loss:
2.3611 - accuracy: 0.1000 - val_loss: 2.3635 - val_accuracy: 0.1000



Run2

Epoch 200/200


196/196 [=====] - 10s 52ms/step - loss:
1.5052 - accuracy: 0.9559 - val_loss: 1.7619 - val_accuracy: 0.6964



Run3

Epoch 200/200

196/196 [=====] - 17s 89ms/step - loss:
1.5708 - accuracy: 0.8900 - val_loss: 1.8016 - val_accuracy: 0.6551




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


196/196 [=====] - 10s 52ms/step - loss:
1.5459 - accuracy: 0.9155 - val_loss: 1.8368 - val_accuracy: 0.6260



Run2

Epoch 200/200


196/196 [=====] - 10s 52ms/step - loss:
1.5240 - accuracy: 0.9376 - val_loss: 1.8262 - val_accuracy: 0.6342



Run3

Epoch 200/200

196/196 [=====] - 22s 112ms/step - loss:
1.5517 - accuracy: 0.9100 - val_loss: 1.8542 - val_accuracy: 0.6023



CIFAR-10: VGG-C


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


196/196 [=====] - 11s 57ms/step - loss:
1.4909 - accuracy: 0.9704 - val_loss: 1.6888 - val_accuracy: 0.7715



Run2

Epoch 200/200


196/196 [=====] - 11s 58ms/step - loss:
1.4935 - accuracy: 0.9676 - val_loss: 1.7032 - val_accuracy: 0.7583



Run3

Epoch 200/200

196/196 [=====] - 11s 56ms/step - loss:
1.4900 - accuracy: 0.9712 - val_loss: 1.6792 - val_accuracy: 0.7796




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


196/196 [=====] - 11s 56ms/step - loss:
1.5204 - accuracy: 0.9418 - val_loss: 1.7004 - val_accuracy: 0.7604



Run2

Epoch 200/200


196/196 [=====] - 11s 56ms/step - loss:
1.5148 - accuracy: 0.9472 - val_loss: 1.6971 - val_accuracy: 0.7653



Run3

Epoch 200/200

196/196 [=====] - 11s 56ms/step - loss:
1.5487 - accuracy: 0.9129 - val_loss: 1.8216 - val_accuracy: 0.6362



CIFAR-10: VGG-D


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


196/196 [=====] - 27s 137ms/step - loss:
1.5072 - accuracy: 0.9540 - val_loss: 1.6962 - val_accuracy: 0.7660



Run2

Epoch 200/200


196/196 [=====] - 13s 65ms/step - loss:
1.4994 - accuracy: 0.9618 - val_loss: 1.6756 - val_accuracy: 0.7844



Run3

Epoch 200/200

196/196 [=====] - 22s 110ms/step - loss:
1.5015 - accuracy: 0.9597 - val_loss: 1.6795 - val_accuracy: 0.7814




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


196/196 [=====] - 13s 67ms/step - loss:
1.5349 - accuracy: 0.9266 - val_loss: 1.7076 - val_accuracy: 0.7523



Run2

Epoch 200/200


196/196 [=====] - 28s 141ms/step - loss:
1.5149 - accuracy: 0.9465 - val_loss: 1.6859 - val_accuracy: 0.7757



Run3

Epoch 200/200

196/196 [=====] - 27s 138ms/step - loss:
1.5894 - accuracy: 0.8700 - val_loss: 1.7404 - val_accuracy: 0.7193



CIFAR-10: VGG-E


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


196/196 [=====] - 15s 76ms/step - loss:
1.8886 - accuracy: 0.5714 - val_loss: 1.9584 - val_accuracy: 0.4978



Run2

Epoch 200/200


196/196 [=====] - 15s 78ms/step - loss:
2.3612 - accuracy: 0.1000 - val_loss: 2.3591 - val_accuracy: 0.1000



Run3

Epoch 200/200

196/196 [=====] - 15s 76ms/step - loss:
2.3612 - accuracy: 0.1000 - val_loss: 2.3591 - val_accuracy: 0.1000




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


196/196 [=====] - 32s 165ms/step - loss: 1.9047 - accuracy: 0.5552 - val_loss: 1.9682 - val_accuracy: 0.4882



Run2

Epoch 200/200


196/196 [=====] - 15s 76ms/step - loss: 1.9984 - accuracy: 0.4587 - val_loss: 2.0600 - val_accuracy: 0.3983



Run3

Epoch 200/200

196/196 [=====] - 15s 76ms/step - loss: 1.6355 - accuracy: 0.8252 - val_loss: 1.7266 - val_accuracy: 0.7332



CIFAR-10: GoogLeNet


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.4

Run1

Epoch 200/200

50000/50000 [=====] - 10s 193us/sample - loss: 2.3612 - accuracy: 0.1000 - val_loss: 2.3612 - val_accuracy: 0.1000



Run2


Epoch 200/200

196/196 [=====] - 9s 47ms/step - loss: 2.3613 - accuracy: 0.1000 - val_loss: 2.3591 - val_accuracy: 0.1000

Run3

Epoch 200/200

50000/50000 [=====] - 10s 192us/sample - loss: 2.3611 - accuracy: 0.1000 - val_loss: 2.3612 - val_accuracy: 0.1000




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.4

Run1

Epoch 200/200


50000/50000 [=====] - 10s 190us/sample - loss:
1.6240 - accuracy: 0.8368 - val_loss: 1.7830 - val_accuracy: 0.6774



Run2

Epoch 200/200


50000/50000 [=====] - 10s 193us/sample - loss:
1.6696 - accuracy: 0.7914 - val_loss: 1.8142 - val_accuracy: 0.6454



Run3

Epoch 200/200

50000/50000 [=====] - 10s 192us/sample - loss:
1.6335 - accuracy: 0.8272 - val_loss: 1.8009 - val_accuracy: 0.6591



MNIST: VGG-A


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


60000/60000 [=====] - 9s 147us/sample - loss:
1.4623 - accuracy: 0.9988 - val_loss: 1.4693 - val_accuracy: 0.9919



Run2

Epoch 200/200


235/235 [=====] - 9s 36ms/step - loss:
1.4623 - accuracy: 0.9988 - val_loss: 1.4692 - val_accuracy: 0.9917



Run3

Epoch 200/200

60000/60000 [=====] - 9s 147us/sample - loss:
1.4621 - accuracy: 0.9991 - val_loss: 1.4697 - val_accuracy: 0.9913




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


60000/60000 [=====] - 9s 148us/sample - loss:
1.4626 - accuracy: 0.9986 - val_loss: 1.4698 - val_accuracy: 0.9914



Run2

Epoch 200/200


60000/60000 [=====] - 9s 147us/sample - loss:
1.4625 - accuracy: 0.9986 - val_loss: 1.4705 - val_accuracy: 0.9905



Run3

Epoch 200/200

60000/60000 [=====] - 9s 147us/sample - loss:
1.4627 - accuracy: 0.9984 - val_loss: 1.4713 - val_accuracy: 0.9896



MNIST: VGG-B


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


235/235 [=====] - 11s 46ms/step - loss:
1.4622 - accuracy: 0.9989 - val_loss: 1.4693 - val_accuracy: 0.9916



Run2

Epoch 200/200


60000/60000 [=====] - 11s 185us/sample - loss:
1.4622 - accuracy: 0.9989 - val_loss: 1.4700 - val_accuracy: 0.9912



Run3

Epoch 200/200

60000/60000 [=====] - 11s 190us/sample - loss:
1.4622 - accuracy: 0.9990 - val_loss: 1.4691 - val_accuracy: 0.9921




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


235/235 [=====] - 39s 166ms/step - loss:
1.4624 - accuracy: 0.9987 - val_loss: 1.4689 - val_accuracy: 0.9922



Run2

Epoch 200/200


235/235 [=====] - 37s 157ms/step - loss:
1.4626 - accuracy: 0.9985 - val_loss: 1.4710 - val_accuracy: 0.9902



Run3

Epoch 200/200

235/235 [=====] - 11s 46ms/step - loss:
1.4625 - accuracy: 0.9987 - val_loss: 1.4710 - val_accuracy: 0.9901



MNIST: VGG-C


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


60000/60000 [=====] - 12s 205us/sample - loss:
1.4624 - accuracy: 0.9988 - val_loss: 1.4709 - val_accuracy: 0.9906



Run2

Epoch 200/200


235/235 [=====] - 12s 50ms/step - loss:
1.4621 - accuracy: 0.9991 - val_loss: 1.4670 - val_accuracy: 0.9940



Run3

Epoch 200/200

60000/60000 [=====] - 12s 201us/sample - loss:
1.4626 - accuracy: 0.9986 - val_loss: 1.4693 - val_accuracy: 0.9918




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


60000/60000 [=====] - 12s 200us/sample - loss:
1.4621 - accuracy: 0.9991 - val_loss: 1.4683 - val_accuracy: 0.9926



Run2

Epoch 200/200


235/235 [=====] - 12s 49ms/step - loss:
1.4627 - accuracy: 0.9985 - val_loss: 1.4730 - val_accuracy: 0.9878



Run3

Epoch 200/200

60000/60000 [=====] - 12s 202us/sample - loss:
1.4628 - accuracy: 0.9984 - val_loss: 1.4721 - val_accuracy: 0.9891



MNIST: VGG-D


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


235/235 [=====] - 14s 58ms/step - loss:
1.4623 - accuracy: 0.9989 - val_loss: 1.4688 - val_accuracy: 0.9919



Run2

Epoch 200/200


60000/60000 [=====] - 14s 241us/sample - loss:
1.4624 - accuracy: 0.9987 - val_loss: 1.4710 - val_accuracy: 0.9903



Run3

Epoch 200/200

60000/60000 [=====] - 15s 248us/sample - loss:
1.4622 - accuracy: 0.9990 - val_loss: 1.4679 - val_accuracy: 0.9930




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


60000/60000 [=====] - 14s 236us/sample - loss:
1.4625 - accuracy: 0.9986 - val_loss: 1.4701 - val_accuracy: 0.9914



Run2

Epoch 200/200


60000/60000 [=====] - 14s 240us/sample - loss:
1.4626 - accuracy: 0.9986 - val_loss: 1.4695 - val_accuracy: 0.9915



Run3

Epoch 200/200

235/235 [=====] - 14s 61ms/step - loss:
1.4625 - accuracy: 0.9987 - val_loss: 1.4703 - val_accuracy: 0.9909



MNIST: VGG-E


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


235/235 [=====] - 17s 70ms/step - loss:
1.4622 - accuracy: 0.9989 - val_loss: 1.4687 - val_accuracy: 0.9920



Run2

Epoch 200/200


60000/60000 [=====] - 17s 283us/sample - loss:
1.4627 - accuracy: 0.9985 - val_loss: 1.4706 - val_accuracy: 0.9905



Run3

Epoch 200/200

60000/60000 [=====] - 17s 283us/sample - loss:
1.4624 - accuracy: 0.9987 - val_loss: 1.4692 - val_accuracy: 0.9920




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.005, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


235/235 [=====] - 17s 71ms/step - loss:
1.4680 - accuracy: 0.9933 - val_loss: 1.4722 - val_accuracy: 0.9886



Run2

Epoch 200/200


235/235 [=====] - 17s 71ms/step - loss:
1.4635 - accuracy: 0.9977 - val_loss: 1.4673 - val_accuracy: 0.9937



Run3

Epoch 200/200

235/235 [=====] - 60s 253ms/step - loss:
1.4644 - accuracy: 0.9967 - val_loss: 1.4697 - val_accuracy: 0.9910



MNIST: GoogLeNet


LR 0.01

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


60000/60000 [=====] - 11s 184us/sample - loss:
1.4623 - accuracy: 0.9989 - val_loss: 1.4676 - val_accuracy: 0.9935



Run2

Epoch 200/200


60000/60000 [=====] - 11s 181us/sample - loss:
1.4622 - accuracy: 0.9990 - val_loss: 1.4675 - val_accuracy: 0.9936



Run3

Epoch 200/200

60000/60000 [=====] - 11s 180us/sample - loss:
1.4622 - accuracy: 0.9990 - val_loss: 1.4681 - val_accuracy: 0.9932




LR 0.005

Optimizer=SGD, Epochs=200, Learning_Rate=0.01, Batch_Size=256, Momentum=0.9
Dropout=0.5

Run1

Epoch 200/200


60000/60000 [=====] - 11s 181us/sample - loss:
1.4628 - accuracy: 0.9983 - val_loss: 1.4697 - val_accuracy: 0.9913



Run2

Epoch 200/200

60000/60000 [=====] - 11s 180us/sample - loss:
1.5549 - accuracy: 0.9012 - val_loss: 1.5622 - val_accuracy: 0.8940



Run3

Epoch 200/200

60000/60000 [=====] - 11s 179us/sample - loss:
1.4628 - accuracy: 0.9984 - val_loss: 1.4687 - val_accuracy: 0.9925

